

Week 1

① supervised learning
"right answers" given

② Classification - Discrete valued output
Regression - Continuous valued output

③ Unsupervised Learning - Let the program find structure
clustering

Week 2

① $x_j^{(i)}$ = value of feature j in the i^{th} training example

$$x_0^{(i)} = 1 \quad \theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

② For gradient descent, feature scaling can make it much faster.

Get every feature into approximately a $-1 \leq x_i \leq 1$ range

Mean normalization: make features have approximately 0 mean

$$x_i \leftarrow \frac{x_i - \mu_i}{s_i}$$

μ_i : average

s_i : range or standard deviation

③ Normal Equation: $\theta = (X^T X)^{-1} X^T y$

Octave: `pinv(X' * X) * X' * y`

no need of feature scaling

pro: ① 不用选 α

② 不用 iterate

$$X = \begin{bmatrix} 1 & x_1 & x_2 & x_3 & x_4 & \dots \\ & \vdots & & & & \end{bmatrix} \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \end{bmatrix}$$

con: slow when n is large (10^4)

因为算 inverse matrix 很慢

If not invertible: ① redundant feature ② too many features

④ Octave

- a. 不等: \sim
- b. comment: %
- c. disp()
- d. eye(4) 出来 4x4 identity matrix
- e. size() length()
- f. who 显示内存里的变量
whos 更详细的信息
- g. clear ... 删变量
- h. save hello.mat v 存 v 进 hello.mat (binary form)
save hello.txt v -ascii (ascii form)
- i. A(:) put all elements of A into a single vector
- j. sum() floor() ceil() prod()
- k. max(A, [], 1) per column
max(A, [], 2) per row
- l. print -dpng 'myplot.png' 把 plot 出来的存成 png

for loop:

```
for i = 1:10,  
    V(i) = 2 ^ i;  
end;
```

while loop:

```
while i <= 5,  
    V(i) = 100;  
    i = i + 1;  
end;
```

define function 在 .m 文件里

```
function y = square This Number(x)  
y = x ^ 2;
```

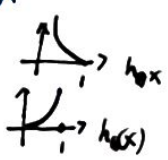
```
function [y1, y2] = choose two (x)
```

⑤ Logistic Regression Model - Want $0 \leq h_{\theta}(x) \leq 1$ It is convex.

$$h_{\theta}(x) = g(\theta^T x) \text{ where } g(z) = \frac{1}{1+e^{-z}} \quad (\text{sigmoid/logistic function})$$

$h_{\theta}(x)$ = estimated probability that $y=1$ on input x

$$\text{cost function, } \text{cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y=1 \\ -\log(1-h_{\theta}(x)) & \text{if } y=0 \end{cases}$$



$$= -y \log(h_{\theta}(x)) - (1-y) \log(1-h_{\theta}(x))$$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

$$= -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1-y^{(i)}) \log (1-h_{\theta}(x^{(i)})) \right]$$

$$\frac{\partial}{\partial \theta} J(\theta) = \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \cdot \frac{1}{n}$$

$$\theta = \theta - \frac{\alpha}{m} X^T (y(X\theta) - \vec{y})$$

advanced algorithm:

① Conjugate gradient

② BFGS

③ L-BFGS

way to implement advanced algorithms:

function [jVal, gradient] = costFunction(theta)

jVal = [... code to compute J(θ)]

gradient = [... code to compute gradient]

end

options = optimset('GradObj', 'On', 'MaxIter', '100');

initialTheta = zeros(2,1)

[optTheta, functionVal, exitFlag] = fminunc(@costFunction, initialTheta, options);

bool: 是否收敛

printer

☆ 当有多种分类时: multiclass classification: one-vs-all

⑥ Overfitting Problem

Solution: i. - 筛选 which features to keep
- Model selection algorithm

ii. Regularization

- keep all features but reduce magnitude of parameters
- 当有很多 features 而它们全有用的时候用

$$\min \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2$$

gradient
descent

$$\Rightarrow \theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j (1 - \alpha \frac{\lambda}{m}) - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

normal
equation

$$\Rightarrow \theta = (X^T X + \lambda A)^{-1} X^T y \quad A = \begin{bmatrix} 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & 0 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & \dots & 0 & 1 \end{bmatrix}$$

For logistic regression

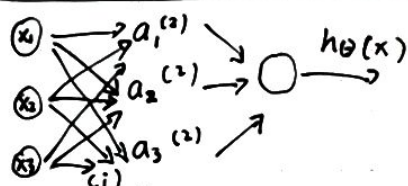
$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \boxed{\frac{\lambda}{m} \theta_j} \quad \text{当 } j=0, \text{ 蓝框项为 } 0$$

Week 4

① Neural Network

If network has s_j units in layer j
 s_{j+1} units in layer $j+1$, then $\theta^{(j)}$ will
be of dimension $s_{j+1} \times (s_j + 1)$



a_i "activation" of unit i in layer j

$\theta^{(j)}$ matrix of weights controlling
function mapping from layer
 j to $j+1$

Week 5

① Neural Network (Classification)

L : total # of layers in network

S_L : # of layers in layer l

Cost Function: $h_\theta(x) \in R^k$ ($h_\theta(x)_i$ = i th output)

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K \left[y_k^{(i)} \log(h_\theta(x^{(i)}))_k + (1 - y_k^{(i)}) \log(1 - (h_\theta(x^{(i)}))_k) \right] + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{S_{l+1}} \sum_{j=1}^{S_l} (\theta_{ji}^{(l)})^2$$

- The double sum simply adds up the logistic regression costs calculated for each cell in the output layer.
- The triple sum adds up squares of all the individual θ s in the entire network.
- The i in the triple sum does not refer to training example i .

In order to get $\frac{\partial}{\partial \theta_{ij}^{(l)}} J(\theta)$ to minimize $J(\theta)$, we use Backpropagation algorithm

With training set $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

a. Set $\Delta_{ij}^{(L)} = 0$ (for all i, j, l)

b. For $i = 1$ to m

i. set $a^{(1)} = x^{(i)}$

ii. perform forward propagation to compute $a^{(l)}$ for $l = 2, 3, \dots, L$

iii. Using $y^{(i)}$ compute $\delta^{(L)} = a^{(L)} - y^{(i)}$

Use BP \leftarrow iv. Compute $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$

v. $\Delta_{ij}^{(L)} = \delta^{(L+1)} (a^{(L)})^T$

$\Delta_{ij}^{(l)} = a_j^{(l)} \delta_i^{(l+1)}$

c. $D_{ij}^{(l)} = \frac{1}{m} \Delta_{ij}^{(l)} + \lambda \theta_{ij}^{(l)}$ if $j \neq 0$

$D_{ij}^{(l)} = \frac{1}{m} \Delta_{ij}^{(l)}$ if $j = 0$

BP:

$$\delta^{(l)} = ((\theta^{(l+1)})^T \delta^{(l+1)}) \cdot * a^{(l)} \cdot * (1 - a^{(l)})$$

$$\frac{\partial}{\partial \theta_{ij}^{(l)}} J(\theta) = D_{ij}^{(l)}$$

为了用 fminunc, 传 theta 前要用 theta(:) unroll, cost function 传出来的 gradient 也要 unroll. 重组可以用 reshape 从 vector 变回 matrix

We can use gradient checking to verify we are getting right gradient.

$$\frac{\partial}{\partial \theta_j} J(\theta) \approx \frac{J(\theta_1, \dots, \theta_j + \epsilon, \dots, \theta_n) - J(\theta_1, \dots, \theta_j - \epsilon, \dots, \theta_n)}{2\epsilon}$$

神经网络的 θ 需要 random initialization

$$\theta = \text{rand}(m, n) * (2 * \text{INIT_EPSILON}) - \text{INIT_EPSILON}$$

Week 6

① Evaluating a Hypothesis

fixing We adjust the algorithm by:

high variance getting more training example

high variance Trying smaller sets of features

bias • Trying additional features

• Increasing or decreasing λ
variance bias

$$\text{err}(h_\theta(x), y) = \begin{cases} 1 & \text{if } h_\theta(x) \geq 0.5 \text{ and } y=0 \text{ or } h_\theta(x) < 0.5 \text{ and } y=1 \\ 0 & \text{otherwise} \end{cases}$$

a. Learn θ using training set

b. compute test error

Linear Regression:

$$J_{\text{test}}(\theta) = \frac{1}{2m_{\text{test}}} \sum_{i=1}^{n_{\text{test}}} (h_\theta(x_{\text{test}}^{(i)}) - y_{\text{test}}^{(i)})^2$$

Classification:

$$\text{test error} = \frac{1}{m_{\text{test}}} \sum_{i=1}^{n_{\text{test}}} \text{err}(h_\theta(x_{\text{test}}^{(i)}), y_{\text{test}}^{(i)})$$

Bias (underfit)

High $J_{\text{train}}(\theta)$ and $J_{\text{cv}}(\theta)$

$$J_{\text{cv}}(\theta) \approx J_{\text{train}}(\theta)$$

Variance (overfit)

High $J_{\text{cv}}(\theta)$ Low $J_{\text{train}}(\theta)$

$$J_{\text{cv}}(\theta) \gg J_{\text{train}}(\theta)$$

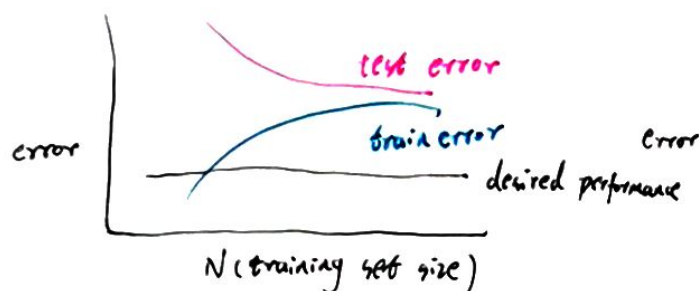
Choosing λ

用不同 λ 和 degree 和 variant 组合学 θ

学的时候用 regularized cost

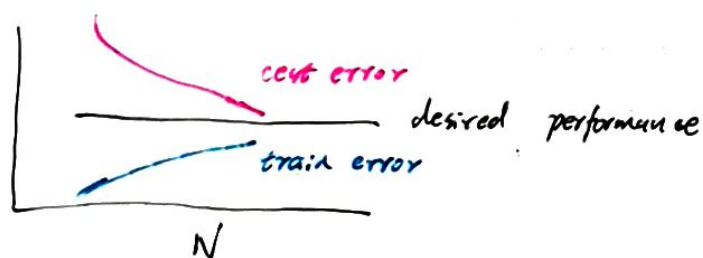
然后算 $J_{\text{cv}}(\theta)$, 选最低 $J_{\text{cv}}(\theta)$

$$\lambda = \{0.01, 0.02, \dots, 10.29\}$$



High bias

Feeding more data does not help



High variance

Feeding more data helps

$$\text{Precision} = \frac{\text{True Positives}}{\# \text{ predicted as positive}}$$

$$\text{Recall} = \frac{\text{True Positives}}{\# \text{ actual positives}}$$

Positive: in presence of rare class that we want to detect

$$F_1 \text{ Score} = 2 \frac{PR}{P+R}$$

Week 7 SVM (convex)

① Kernels

$$h_\theta(x) = \begin{cases} 1 & \text{if } \theta_0 + \theta_1 f_1 + \dots \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

这个 f 是 kernel

都要符合 Mercer's Theorem

polynomial, string, chi-square, histogram intersection

Gaussian kernel: Normalize first

分子是 $(^{(u)})$ 到 x 的距离平方

$$f_1 = \text{similarity}(x, l^{(u)}) = \exp\left(-\frac{\|x - l^{(u)}\|^2}{2\sigma^2}\right) = \exp\left(-\frac{\sum_{j=1}^n (x_j - l_j^{(u)})^2}{2\sigma^2}\right)$$

如果 $x \approx l^{(u)}$: $f_1 \approx 1$

如果 x is far from $l^{(u)}$: $f_1 \approx 0$

σ^2 越小, 对距离要求更严格 lower bias, higher variance

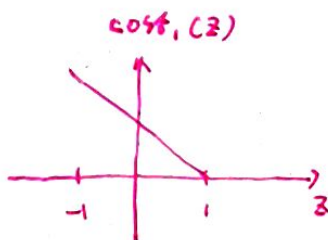
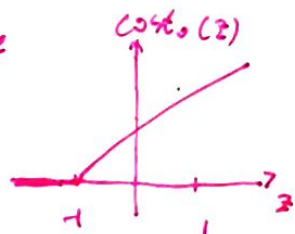
什么时候用 Logistic Regression or SVM?

- ① 当 n is large (relative to m): 用 Linear Regression or SVM without kernel
- ② 当 n is small, m is intermediate: SVM with Gaussian Kernel
- ③ 当 n is small, m is large: add more features
then use logistic regression or SVM without kernel

The SVM solves

$$\min_{\theta} C \sum_{i=1}^n \left[y^{(i)} \text{cost}_1(\theta^T f^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T f) \right] + \sum_{j=1}^n \theta_j^2 \cdot \frac{1}{2}$$

where



classification problem
provide maximum margin

如果 $C = \frac{1}{\lambda}$, 得出的答案跟 logistic regression 一样, C 通常很大

★ SVM VS Logistic Regression

- ① Logistic Regression is more sensitive to outlier.
- ② LR 给 probability, SVM 直接给 0 或 1 的结果

结论

先用 LR, 如果发现不是 linearly separable 用 SVM with non-linear kernel

SVM software package: liblinear, libsvm ... 都要符合 Mercer's Theorem

Need to specify: ① Choice of parameter C

② Choice of kernel

很多 packages 已有内置 multi-class classification, 没有的话用 one-vs-all

Week 8 Unsupervised Machine Learning + PCA

K-Means Method

- ① $c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is assigned
 μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$)
 $\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

② K-Means Objective

$$\min_{c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \min \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

③ K-Means Algorithm

- a. 随机 initialize K cluster centroids $\mu_1, \dots, \mu_K \in \mathbb{R}^n$
在 1 到 m 随机选 K 个数然后以那些数在 x 里选点做 μ
有机会 stuck 在 local minimum, 要多用几次不同随机起点找最低 J
- b. Repeat {
 for $i=1$ to m
 $c^{(i)} := \text{index (from 1 to } K) \text{ of cluster centroid closest to } x^{(i)}$
 for $k=1$ to K
 $\mu_k := \text{mean of points assigned to cluster } k$
}

④ 选 K 的方法 (怎么知道选多少个 cluster)

Elbow Method: 一个个去试, 找 J 转折点
不过有时候没转折点

⑤ Random Initialization

$$\text{randIdx} = \text{randperm}(\text{size}(X, 1));$$

$$\text{centroids} = X(\text{randIdx}(1:K), :);$$

Principal Component Analysis

用途:

1. Compression: 从 n -dimensions 压成 k -dimensions

① Reduce memory / disk needed to store data

② speed up learning algorithm

当真的慢或者真没内存再用 PCA, 不是首选的优化方法

2. Visualization: $k=2$ or 3

不应该使用的情况: To prevent overfitting, 应该用 regularization

算法 Algorithm

① 预处理: mean normalization / feature scaling

② 计算 "covariance matrix"

$$\Sigma = \frac{1}{m} \sum_{i=1}^m (x^{(i)})(x^{(i)})^T$$

vectorized:

$$\text{Sigma} = (1/m) * X^T * X$$

③ 计算 "eigenvectors" of matrix Σ

$$[U, S, V] = \text{svd}(\text{Sigma});$$

$$U = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \dots & u^{(n)} \\ | & | & & | \end{bmatrix} \quad S = \begin{bmatrix} s_{11} & & 0 \\ & s_{22} & \\ 0 & & \ddots \\ & & & s_{nn} \\ & & 0 & & \end{bmatrix}$$

④ 得出新坐标 Z

压缩 dimension: $U = U(:, 1:k)$

vectorized:

$$Z = U^T X$$

$$Z = X * U$$

Reconstruction: $X_{\text{approx}}^{(i)} = U_{\text{reduced}} \cdot Z^{(i)}$

$$X_{\text{rec}} = Z * U_{\text{reduced}}^T$$

怎么选 k ? (压缩到什么 dimension)

Typically, choose k to be smallest value so that

$$\frac{\sum_{i=1}^m \|x^{(i)} - x_{\text{approx}}^{(i)}\|^2}{\sum_{i=1}^m \|x^{(i)}\|^2} \leq 0.01$$

$\Rightarrow 99\%$ of variance is retained

用 svd: for given k :

$$1 - \frac{\sum_{i=1}^k s_{ii}}{\sum_{i=1}^n s_{ii}} \leq 0.01$$

Week 9 Anomaly Detection + Recommender Systems

(1)

Anomaly Detection Algorithm

1. Choose feature x_i that we think might be indicative of anomalous examples

2. Calculate parameters $\mu_1, \dots, \mu_n, \sigma_1^2, \dots, \sigma_n^2$

$$\mu_j = \frac{1}{m} \sum_{i=1}^m x_j^{(i)}$$

$m \uparrow$ sample $n \uparrow$ feature

$$\sigma_j^2 = \frac{1}{m} \sum_{i=1}^m (x_j^{(i)} - \mu_j)^2$$

3. Given new example x , compute $p(x)$

$$p(x) = \prod_{j=1}^n P(x_j; \mu_j, \sigma_j^2)$$

$$= \prod_{j=1}^n \frac{1}{\sqrt{2\pi}\sigma_j} \exp\left(-\frac{(x_j - \mu_j)^2}{2\sigma_j^2}\right) \quad (\text{Normal Distribution})$$

Anomaly if $p(x) < \epsilon$

Training set 全是正常的, 用来学习 μ_j 和 σ_j^2

Validation set 混合正常 + 不正常, 用来调整 ϵ

Test set 混合正常 + 不正常, 用 F1 score 评价表现

什么时候用 Anomaly Detection

or

Supervised Learning

① 很少 positive example

② 很多 negative example

③ 异常的种类千奇百怪

很可能跟已收集到的异常不同

有许多 positive 和 negative 的例子

可供学习

例子

① Fraud Detection

② Manufacturing

③ Monitoring machines at a data center

① Email spam detection

② weather prediction

③ Cancer classification

注意：要把 non-Gaussian feature 转换为 Gaussian

考虑的方法： $\log(x_i + c)$

$$x_i \wedge c$$

Original Model

① 要自己创造 feature 来考虑

x_i, x_j 之间的关系在内

② 资源要求小，适合多 feature

13. Multivariate Gaussian

① 自动考虑 feature 间的关系

② 资源要求大

③ 一定要 $m \gg n$

Anomaly Detection with multivariate Gaussian Algorithm

1. Fit model $p(x)$ by setting

$$\mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

$$\Sigma = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu)(x^{(i)} - \mu)^T$$

2. Given a new example x , compute

$$p(x) = \frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma|^{\frac{1}{2}}} \exp\left(-\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu)\right)$$

flag anomaly if $p(x) < \epsilon$

作业中问题：

① $(A * R) \wedge 2$ 不加蓝括号会失败

② logical array 改里面数字要重定义为 float 或 int 8

Collaborative Filtering Algorithm

1. Initialize $x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_m)}$ to small random values ^{to break symmetry}
2. Minimize $J(x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_m)})$ using gradient descent (or an advanced optimization algorithm)

$$J = \frac{1}{2} \sum_{(i,j): r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

$$x_k^{(i)} = x_k^{(i)} - \alpha \left(\sum_{j: r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) \theta_k^{(j)} + \lambda x_k^{(i)} \right) = \frac{\partial}{\partial x_k^{(i)}} J$$

$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \left(\sum_{i: r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) x_k^{(i)} + \lambda \theta_k^{(j)} \right) = \frac{\partial}{\partial \theta_k^{(j)}} J$$

$r(i,j) = 1$ if user j has rated movie i (0 otherwise)

$y^{(i,j)}$ = rating by user j on movie i

$\theta^{(j)}$ = parameter vector for user j

$x^{(i)}$ = feature vector for movie i

For user j , movie i , predicted rating: $(\theta^{(j)})^T x^{(i)}$

$$X = \begin{bmatrix} -(x^{(1)})^T & - \\ \vdots & \\ -(x^{(n_m)})^T & - \end{bmatrix} \quad \Theta = \begin{bmatrix} -(\theta^{(1)})^T & - \\ \vdots & \\ -(\theta^{(n_u)})^T & - \end{bmatrix} \quad \text{vectorized:} \quad Y = X \Theta^T$$

算 cost 和 gradient 的 MATLAB vectorized 代码

```
estimated_error = ((X * Theta' - Y) .* R);
```

```
J = 0.5 * sum(estimated_error.^2, 'all');
```

```
J = J + lambda/2 * (sum(Theta.^2, 'all') + sum(X.^2, 'all'));
```

```
for i = 1:num_movies
```

```
    for k = 1:num_features
```

```
        X_grad(i,k) = sum(estimated_error(i,:)'.* Theta(:,k));
```

```
    end
```

```
    X_grad(i,:) = X_grad(i,:) + lambda * X(i,:);
```

```
end
```

```
for j = 1:num_users
```

```
    for k = 1:num_features
```

```
        Theta_grad(j,k) = sum(estimated_error(:,j).* X(:,k));
```

```
    end
```

```
    Theta_grad(j,:) = Theta_grad(j,:) + lambda * Theta(j,:);
```

```
end
```

Week 10

① Stochastic Gradient Descent Algorithm

每次(步)根据一个 sample 优化全局参数 比 batch 快

1. Randomly shuffle (reorder) training examples

2. Repeat {

for $i = 1, \dots, m$ {

$$\theta_j := \theta_j - \alpha (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \quad (\text{for every } j = 0, \dots, n)$$

}

}

② Mini-batch gradient descent

1. Shuffle

Say $b = 10, m = 1000$

2. Repeat {

for $i = 1, 11, 21, 31, \dots, 991$ {

$$\theta_j := \theta_j - \alpha \frac{1}{b} \sum_{k=i}^{i+b-1} (h_{\theta}(x^{(k)}) - y^{(k)}) x_j^{(k)} \quad (\text{for every } j = 0, \dots, n)$$

}

}

①和②都可以动态变化学习率 α , 越往后越小来求更好的参数

③ Online learning

适合 continuous data flow

来一个优化一次然后舍弃

④ Map Reduce and Data parallelism

只要算法是 training set 的某种和就可分拆平行计算

① Pipeline : 流程.

机器学习中经常把一件工作分拆成几个工序

比如 Photo OCR: Image \rightarrow Text Detection \rightarrow Character Segmentation
 \rightarrow Character Recognition

② Artificial Data Synthesis

可以通过自己对原始数据集的加工来扩大训练集

比如: 清晰语音加背景噪音

图像的变形

③ Ceiling Analysis

用来分析 pipeline 中哪些工序更需要改进

人工标记出输出那个工序的完美结果, 然后观察整体算法性能提升幅度